

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>

Using Loss Functions

```
model.compile(loss='mse', optimizer='sgd')
```

or

```
from tensorflow.keras.losses import mean_squared_error  
model.compile(loss=mean_squared_error, optimizer='sgd')
```

Using Loss Functions

```
model.compile(loss='mse', optimizer='sgd')
```

or

```
from tensorflow.keras.losses import mean_squared_error  
model.compile(loss=mean_squared_error, optimizer='sgd')
```

Using Loss Functions

```
model.compile(loss='mse', optimizer='sgd')
```

or

```
from tensorflow.keras.losses import mean_squared_error  
model.compile(loss=mean_squared_error, optimizer='sgd')
```

Using Loss Functions

```
model.compile(loss='mse', optimizer='sgd')
```

or

```
from tensorflow.keras.losses import mean_squared_error  
model.compile(loss=mean_squared_error(param=value),  
optimizer='sgd')
```

Creating a custom loss function

```
def my_loss_function(y_true, y_pred):
```

Creating a custom loss function

```
def my_loss_function(y_true, y_pred):
```

Creating a custom loss function

```
def my_loss_function(y_true, y_pred):
```



Creating a custom loss function

```
def my_loss_function(y_true, y_pred):  
    return losses
```

Example

Huber Loss


$$L_{\delta}(a) \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta \\ \delta \times (|a| - \frac{1}{2}\delta) & \text{otherwise} \end{cases}$$



https://en.wikipedia.org/wiki/Huber_loss

$$L_{\delta}(a) \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta \\ \delta \times (|a| - \frac{1}{2}\delta) & \text{otherwise} \end{cases}$$


Threshold



$$L_{\delta}(a) \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta \\ \delta \times (|a| - \frac{1}{2}\delta) & \text{otherwise} \end{cases}$$

Threshold


Error




$$L_{\delta}(a) \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta \\ \delta \times (|a| - \frac{1}{2}\delta) & \text{otherwise} \end{cases}$$

Threshold


Error
(label-prediction)



$$L_{\delta}(a) \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta \\ \delta \times (|a| - \frac{1}{2}\delta) & \text{otherwise} \end{cases}$$



$$L_{\delta}(a) \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta \\ \delta \times (|a| - \frac{1}{2}\delta) & \text{otherwise} \end{cases}$$




```
def my_huber_loss(y_true, y_pred):  
    threshold = 1  
    error = y_true - y_pred  
    is_small_error = tf.abs(error) <= threshold  
    small_error_loss = tf.square(error) / 2  
    big_error_loss = threshold * (tf.abs(error) - (0.5 * threshold))  
    return tf.where(is_small_error, small_error_loss, big_error_loss)
```

$$L_{\delta}(a) \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta \\ \delta \times (|a| - \frac{1}{2}\delta) & \text{otherwise} \end{cases}$$

```
def my_huber_loss(y_true, y_pred):  
    threshold = 1  
    error = y_true - y_pred  
    is_small_error = tf.abs(error) <= threshold  
    small_error_loss = tf.square(error) / 2  
    big_error_loss = threshold * (tf.abs(error) - (0.5 * threshold))  
    return tf.where(is_small_error, small_error_loss, big_error_loss)
```

$$L_{\delta}(a) \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta \\ \delta \times (|a| - \frac{1}{2}\delta) & \text{otherwise} \end{cases}$$

```
def my_huber_loss(y_true, y_pred):  
    threshold = 1  
    error = y_true - y_pred  
    is_small_error = tf.abs(error) <= threshold  
    small_error_loss = tf.square(error) / 2  
    big_error_loss = threshold * (tf.abs(error) - (0.5 * threshold))  
    return tf.where(is_small_error, small_error_loss, big_error_loss)
```

$$L_{\delta}(a) \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta \\ \delta \times (|a| - \frac{1}{2}\delta) & \text{otherwise} \end{cases}$$

```
def my_huber_loss(y_true, y_pred):  
    threshold = 1  
    error = y_true - y_pred  
    is_small_error = tf.abs(error) <= threshold  
    small_error_loss = tf.square(error) / 2  
    big_error_loss = threshold * (tf.abs(error) - (0.5 * threshold))  
    return tf.where(is_small_error, small_error_loss, big_error_loss)
```

$$L_{\delta}(a) \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta \\ \delta \times (|a| - \frac{1}{2}\delta) & \text{otherwise} \end{cases}$$

```
def my_huber_loss(y_true, y_pred):  
    threshold = 1  
    error = y_true - y_pred  
    is_small_error = tf.abs(error) <= threshold  
    small_error_loss = tf.square(error) / 2  
    big_error_loss = threshold * (tf.abs(error) - (0.5 * threshold))  
    return tf.where(is_small_error, small_error_loss, big_error_loss)
```

$$L_{\delta}(a) \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta \\ \delta \times (|a| - \frac{1}{2}\delta) & \text{otherwise} \end{cases}$$

```
def my_huber_loss(y_true, y_pred):  
    threshold = 1  
    error = y_true - y_pred  
    is_small_error = tf.abs(error) <= threshold  
    small_error_loss = tf.square(error) / 2  
    big_error_loss = threshold * (tf.abs(error) - (0.5 * threshold))  
    return tf.where(is_small_error, small_error_loss, big_error_loss)
```

$$L_{\delta}(a) \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta \\ \delta \times (|a| - \frac{1}{2}\delta) & \text{otherwise} \end{cases}$$

```
def my_huber_loss(y_true, y_pred):
    threshold = 1
    error = y_true - y_pred
    is_small_error = tf.abs(error) <= threshold
    small_error_loss = tf.square(error) / 2
    big_error_loss = threshold * (tf.abs(error) - (0.5 * threshold))
    return tf.where(is_small_error, small_error_loss, big_error_loss)
```

Boolean to check

Value if True

Value if False

$$L_{\delta}(a) \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta \\ \delta \times (|a| - \frac{1}{2}\delta) & \text{otherwise} \end{cases}$$

```
model = tf.keras.Sequential([keras.layers.Dense(units=1, input_shape=[1])])  
model.compile(optimizer='sgd', loss='mean_squared_error')
```



```
model = tf.keras.Sequential([keras.layers.Dense(units=1, input_shape=[1])])  
model.compile(optimizer='sgd', loss='mean_squared_error')
```

```
model = tf.keras.Sequential([keras.layers.Dense(units=1, input_shape=[1])])  
model.compile(optimizer='sgd', loss='my_huber_loss')
```

```
def my_huber_loss(y_true, y_pred):  
    threshold = 1  
    error = y_true - y_pred  
    is_small_error = tf.abs(error) <= threshold  
    small_error_loss = tf.square(error) / 2  
    big_error_loss = threshold * (tf.abs(error) - (0.5 * threshold))  
    return tf.where(is_small_error, small_error_loss, big_error_loss)
```

$$L_{\delta}(a) \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta \\ \delta \times (|a| - \frac{1}{2}\delta) & \text{otherwise} \end{cases}$$

```
def my_huber_loss(y_true, y_pred):  
    threshold = 1  
    error = y_true - y_pred  
    is_small_error = tf.abs(error) <= threshold  
    small_error_loss = tf.square(error) / 2  
    big_error_loss = threshold * (tf.abs(error) - (0.5 * threshold))  
    return tf.where(is_small_error, small_error_loss, big_error_loss)
```

```
def my_huber_loss(y_true, y_pred):  
    threshold = 1  
    error = y_true - y_pred  
    is_small_error = tf.abs(error) <= threshold  
    small_error_loss = tf.square(error) / 2  
    big_error_loss = threshold * (tf.abs(error) - (0.5 * threshold))  
    return tf.where(is_small_error, small_error_loss, big_error_loss)
```

```
def my_huber_loss(y_true, y_pred):  
    threshold = 1  
    error = y_true - y_pred  
    is_small_error = tf.abs(error) <= threshold  
    small_error_loss = tf.square(error) / 2  
    big_error_loss = threshold * (tf.abs(error) - (0.5 * threshold))  
    return tf.where(is_small_error, small_error_loss, big_error_loss)
```

```
def my_huber_loss_with_threshold(threshold):  
    def my_huber_loss(y_true, y_pred):  
        error = y_true - y_pred  
        is_small_error = tf.abs(error) <= threshold  
        small_error_loss = tf.square(error) / 2  
        big_error_loss = threshold * (tf.abs(error) - (0.5 * threshold))  
        return tf.where(is_small_error, small_error_loss, big_error_loss)  
    return my_huber_loss
```

```
def my_huber_loss_with_threshold(threshold):  
    def my_huber_loss(y_true, y_pred):  
        error = y_true - y_pred  
        is_small_error = tf.abs(error) <= threshold  
        small_error_loss = tf.square(error) / 2  
        big_error_loss = threshold * (tf.abs(error) - (0.5 * threshold))  
        return tf.where(is_small_error, small_error_loss, big_error_loss)  
    return my_huber_loss
```



```
def my_huber_loss_with_threshold(threshold):
```

```
    def my_huber_loss(y_true, y_pred):
```

```
        error = y_true - y_pred
```

```
        is_small_error = tf.abs(error) <= threshold
```

```
        small_error_loss = tf.square(error) / 2
```

```
        big_error_loss = threshold * (tf.abs(error) - (0.5 * threshold))
```

```
        return tf.where(is_small_error, small_error_loss, big_error_loss)
```

```
    return my_huber_loss
```

```
def my_huber_loss_with_threshold(threshold):  
    def my_huber_loss(y_true, y_pred):  
        error = y_true - y_pred  
        is_small_error = tf.abs(error) <= threshold  
        small_error_loss = tf.square(error) / 2  
        big_error_loss = threshold * (tf.abs(error) - (0.5 * threshold))  
        return tf.where(is_small_error, small_error_loss, big_error_loss)  
    return my_huber_loss
```

```
model.compile(  
    optimizer='sgd', loss=my_huber_loss_with_threshold(threshold=1))
```

```
model.compile(  
    optimizer='sgd', loss=my_huber_loss_with_threshold(threshold=1))
```

```
from tensorflow.keras.losses import Loss
```

```
class MyHuberLoss(Loss):
```

```
    threshold = 1
```

```
    def __init__(self, threshold):
```

```
        super().__init__()
```

```
        self.threshold = threshold
```

```
    def call(self, y_true, y_pred):
```

```
        error = y_true - y_pred
```

```
        is_small_error = tf.abs(error) <= self.threshold
```

```
        small_error_loss = tf.square(error) / 2
```

```
        big_error_loss = self.threshold * (tf.abs(error) - (0.5 * self.threshold))
```

```
        return tf.where(is_small_error, small_error_loss, big_error_loss)
```

```
from tensorflow.keras.losses import Loss
```

```
class MyHuberLoss(Loss):
```

```
    threshold = 1
```

```
    def __init__(self, threshold):
```

```
        super().__init__()
```

```
        self.threshold = threshold
```

```
    def call(self, y_true, y_pred):
```

```
        error = y_true - y_pred
```

```
        is_small_error = tf.abs(error) <= self.threshold
```

```
        small_error_loss = tf.square(error) / 2
```

```
        big_error_loss = self.threshold * (tf.abs(error) - (0.5 * self.threshold))
```

```
        return tf.where(is_small_error, small_error_loss, big_error_loss)
```

```
from tensorflow.keras.losses import Loss
```

```
class MyHuberLoss(Loss):
```

```
    threshold = 1
```

```
    def __init__(self, threshold):
```

```
        super().__init__()
```

```
        self.threshold = threshold
```

```
    def call(self, y_true, y_pred):
```

```
        error = y_true - y_pred
```

```
        is_small_error = tf.abs(error) <= self.threshold
```

```
        small_error_loss = tf.square(error) / 2
```

```
        big_error_loss = self.threshold * (tf.abs(error) - (0.5 * self.threshold))
```

```
        return tf.where(is_small_error, small_error_loss, big_error_loss)
```

```
from tensorflow.keras.losses import Loss
```

```
class MyHuberLoss(Loss):
```

```
    threshold = 1
```

```
    def __init__(self, threshold):
```

```
        super().__init__()
```

```
        self.threshold = threshold
```

```
    def call(self, y_true, y_pred):
```

```
        error = y_true - y_pred
```

```
        is_small_error = tf.abs(error) <= self.threshold
```

```
        small_error_loss = tf.square(error) / 2
```

```
        big_error_loss = self.threshold * (tf.abs(error) - (0.5 * self.threshold))
```

```
        return tf.where(is_small_error, small_error_loss, big_error_loss)
```



```
from tensorflow.keras.losses import Loss
```

```
class MyHuberLoss(Loss):
```

```
    threshold = 1
```

```
    def __init__(self, threshold):
```

```
        super().__init__()
```

```
        self.threshold = threshold
```

```
    def call(self, y_true, y_pred):
```

```
        error = y_true - y_pred
```

```
        is_small_error = tf.abs(error) <= self.threshold
```

```
        small_error_loss = tf.square(error) / 2
```

```
        big_error_loss = self.threshold * (tf.abs(error) - (0.5 * self.threshold))
```

```
        return tf.where(is_small_error, small_error_loss, big_error_loss)
```

```
from tensorflow.keras.losses import Loss
```

```
class MyHuberLoss(Loss):
```

```
    threshold = 1
```

```
    def __init__(self, threshold):
```

```
        super().__init__()
```

```
        self.threshold = threshold
```

```
    def call(self, y_true, y_pred):
```

```
        error = y_true - y_pred
```

```
        is_small_error = tf.abs(error) <= self.threshold
```

```
        small_error_loss = tf.square(error) / 2
```

```
        big_error_loss = self.threshold * (tf.abs(error) - (0.5 * self.threshold))
```

```
        return tf.where(is_small_error, small_error_loss, big_error_loss)
```

```
from tensorflow.keras.losses import Loss
```

```
class MyHuberLoss(Loss):
```

```
    threshold = 1
```

```
    def __init__(self, threshold):
```

```
        super().__init__()
```

```
        self.threshold = threshold
```

```
    def call(self, y_true, y_pred):
```

```
        error = y_true - y_pred
```

```
        is_small_error = tf.abs(error) <= self.threshold
```

```
        small_error_loss = tf.square(error) / 2
```

```
        big_error_loss = self.threshold * (tf.abs(error) - (0.5 * self.threshold))
```

```
        return tf.where(is_small_error, small_error_loss, big_error_loss)
```

```
from tensorflow.keras.losses import Loss
```

```
class MyHuberLoss(Loss):
```

```
    threshold = 1
```

```
    def __init__(self, threshold):
```

```
        super().__init__()
```

```
        self.threshold = threshold
```

```
    def call(self, y_true, y_pred):
```

```
        error = y_true - y_pred
```

```
        is_small_error = tf.abs(error) <= self.threshold
```

```
        small_error_loss = tf.square(error) / 2
```

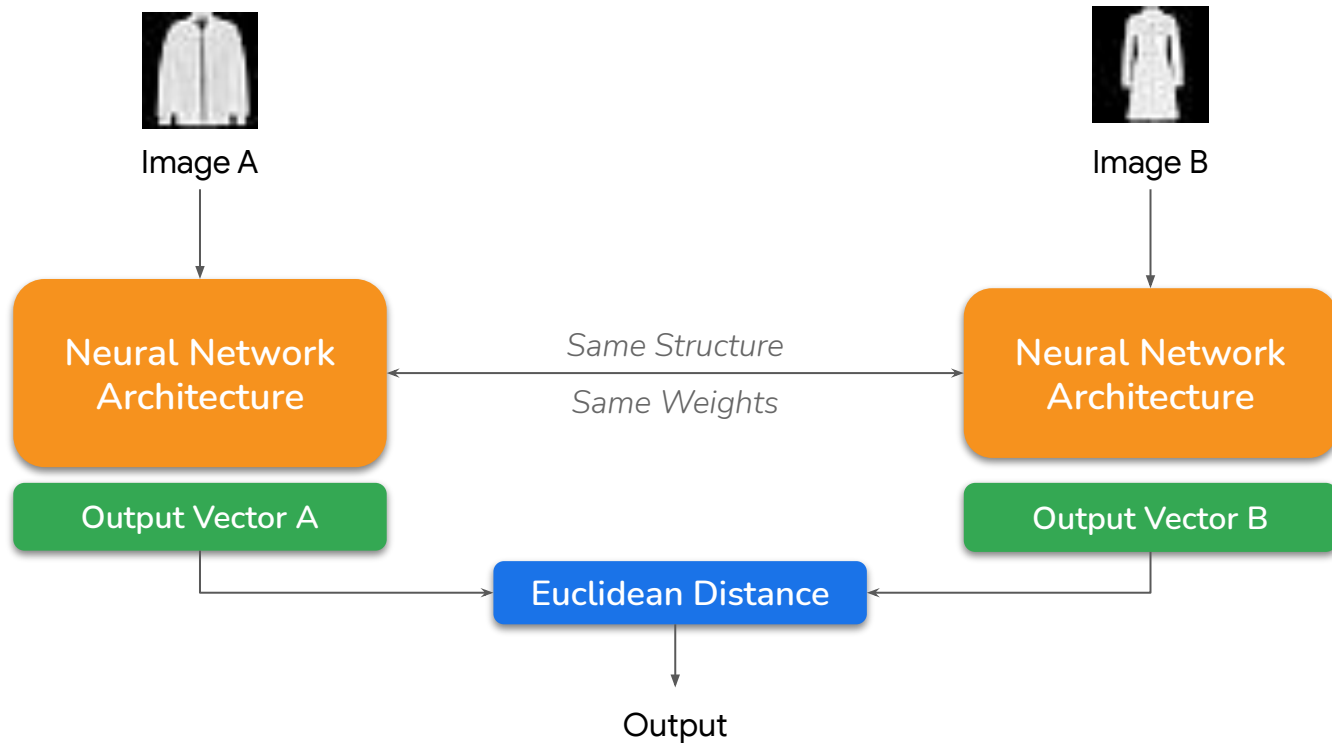
```
        big_error_loss = self.threshold * (tf.abs(error) - (0.5 * self.threshold))
```

```
        return tf.where(is_small_error, small_error_loss, big_error_loss)
```

```
model.compile(optimizer='sgd', loss=MyHuberLoss(threshold=1))
```

```
model.compile(optimizer='sgd', loss=MyHuberLoss(threshold=1))
```

Siamese Network for Image Similarity



Contrastive Loss

- If images are **similar**, produce feature vectors that are **very similar**
- If images are **different**, produce feature vectors that are **dissimilar**.
- Based on the paper

“Dimensionality Reduction by Learning an Invariant Mapping”

by R. Hadsell ; S. Chopra ; Y. LeCun

<http://yann.lecun.com/exdb/publis/pdf/hadsell-chopra-lecun-06.pdf>

Contrastive Loss - Formula

$$Y * D^2 + (1 - Y) * \max(\text{margin} - D, 0)^2$$

Contrastive Loss - Formula

$$\boxed{Y} * D^2 + (1 - \boxed{Y}) * \max(\text{margin} - D, 0)^2$$

Contrastive Loss - Formula


$$Y * D^2 + (1 - Y) * \max(\text{margin} - D, 0)^2$$

Contrastive Loss - Formula


$$Y * D^2 + (1 - Y) * \max(\text{margin} - D, 0)^2$$

Contrastive Loss - Formula

$$Y * D^2 + (1 - Y) * \max(\text{margin} - D, 0)^2$$




$$1 * D^2 + (1 - 1) * \max(\text{margin} - D, 0)^2$$




$$D^2$$

Contrastive Loss - Formula

$$Y * D^2 + (1 - Y) * \max(\text{margin} - D, 0)^2$$



$$0 * D^2 + (1 - 0) * \max(\text{margin} - D, 0)^2$$



$$\max(\text{margin} - D, 0)^2$$


Contrastive Loss - Formula

$$Y * D^2 + (1 - Y) * \max(\text{margin} - D, 0)^2$$

$$Y_{\text{true}} * Y_{\text{pred}}^2 + (1 - Y_{\text{true}}) * \max(\text{margin} - Y_{\text{pred}}, 0)^2$$

Contrastive Loss - Formula

$$Y * D^2 + (1 - Y) * \max(\text{margin} - D, 0)^2$$

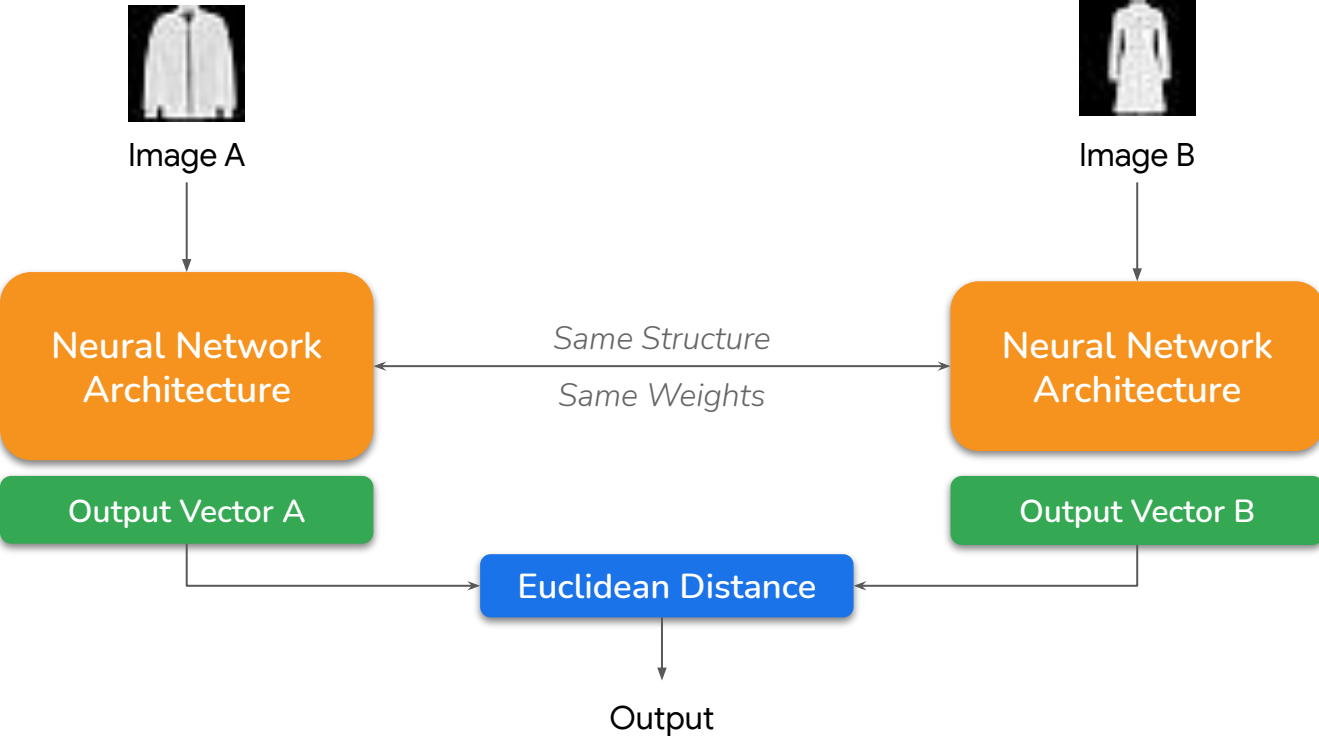

$$Y_{\text{true}} * Y_{\text{pred}}^2 + (1 - Y_{\text{true}}) * \max(\text{margin} - Y_{\text{pred}}, 0)^2$$

Contrastive Loss - Formula

$$Y * D^2 + (1 - Y) * \max(\text{margin} - D, 0)^2$$



$$Y_{true} * Y_{pred}^2 + (1 - Y_{true}) * \max(\text{margin} - Y_{pred}, 0)^2$$



$$Y * D^2 + (1 - Y) * \max(\text{margin} - D, 0)^2$$

Custom Loss Function

```
def contrastive_loss(y_true, y_pred):  
    margin = 1  
    square_pred = K.square(y_pred)  
    margin_square = K.square(K.maximum(margin - y_pred, 0))  
    return K.mean(y_true * square_pred + (1 - y_true) * margin_square)
```

$$Y_{true} * Y_{pred}^2 + (1 - Y_{true}) * \max(\text{margin} - Y_{pred}, 0)^2$$

Custom Loss Function

```
def contrastive_loss(y_true, y_pred):  
    margin = 1  
    square_pred = K.square(y_pred)  
    margin_square = K.square(K.maximum(margin - y_pred, 0))  
    return K.mean(y_true * square_pred + (1 - y_true) * margin_square)
```

$$Y_{true} * Y_{pred}^2 + (1 - Y_{true}) * \max(\text{margin} - Y_{pred}, 0)^2$$

Custom Loss Function

```
def contrastive_loss(y_true, y_pred):  
    margin = 1  
    square_pred = K.square(y_pred)  
    margin_square = K.square(K.maximum(margin - y_pred, 0))  
    return K.mean(y_true * square_pred + (1 - y_true) * margin_square)
```

$$Y_{true} * Y_{pred}^2 + (1 - Y_{true}) * \max(\text{margin} - Y_{pred}, 0)^2$$

Custom Loss Function

```
def contrastive_loss(y_true, y_pred):  
    margin = 1  
    square_pred = K.square(y_pred)  
    margin_square = K.square(K.maximum(margin - y_pred, 0))  
    return K.mean(y_true * square_pred + (1 - y_true) * margin_square)
```

$$K.mean[Y_{true} * Y_{pred}^2 + (1 - Y_{true}) * \max(margin - Y_{pred}, 0)^2]$$

Usage of Custom Loss

```
model.compile(loss=contrastive_loss, optimizer=RMSprop())
```

Usage of Custom Loss

```
model.compile(loss=contrastive_loss, optimizer=RMSprop())
```


Custom Loss Function with Arguments

```
def contrastive_loss_with_margin(margin):
```

```
    #Original Loss Function
```

```
    def contrastive_loss(y_true, y_pred):
```

```
        square_pred = K.square(y_pred)
```

```
        margin_square = K.square(K.maximum(margin - y_pred, 0))
```

```
        return K.mean(y_true * square_pred + (1 - y_true) * margin_square)
```

```
    return contrastive_loss
```

Usage of Wrapper Loss Function

```
model.compile(loss=contrastive_loss_with_margin(margin=1), optimizer=rms)
```

Usage of Wrapper Loss Function

```
model.compile(loss=contrastive_loss_with_margin(margin=1), optimizer=rms)
```

Contrastive Loss - Object Oriented

```
class ContrastiveLoss(Loss):  
    margin = 0  
  
    def __init__(self, margin):  
        super().__init__()  
        self.margin = margin  
  
    def call(self, y_true, y_pred):  
        square_pred = K.square(y_pred)  
        margin_square = K.square(K.maximum(self.margin - y_pred, 0))  
        return K.mean(y_true * square_pred + (1 - y_true) * margin_square)
```

Contrastive Loss - Object Oriented

```
class ContrastiveLoss(Loss):  
    margin = 0  
    def __init__(self, margin):  
        super().__init__()  
        self.margin = margin  
  
    def call(self, y_true, y_pred):  
        square_pred = K.square(y_pred)  
        margin_square = K.square(K.maximum(self.margin - y_pred, 0))  
        return K.mean(y_true * square_pred + (1 - y_true) * margin_square)
```

Contrastive Loss - Object Oriented

```
class ContrastiveLoss(Loss):  
    margin = 0  
  
    def __init__(self, margin):  
        super().__init__()  
        self.margin = margin  
  
    def call(self, y_true, y_pred):  
        square_pred = K.square(y_pred)  
        margin_square = K.square(K.maximum(self.margin - y_pred, 0))  
        return K.mean(y_true * square_pred + (1 - y_true) * margin_square)
```

Usage of Object Oriented Loss

```
model.compile(loss=ContrastiveLoss(margin=1), optimizer=rms)
```

Usage of Object Oriented Loss

```
model.compile(loss=ContrastiveLoss(margin=1), optimizer=rms)
```